



Continuous training with multiple SDKs, Kubeflow, and AI Platform Pipelines



Agenda


Containerized training applications

Containerizing PyTorch, Scikit, and XGBoost applications

Kubeflow and AI Platform Pipelines


Continuous training

Why **use** containerized training applications?




Why use containerized training applications?

- You don't have to worry about dependencies.




Why use containerized training applications?

- You don't have to worry about dependencies.
- Use them as ops in a Kubeflow pipeline (or other orchestration tools).



Why use containerized training applications?

- You don't have to worry about dependencies.
- Use them as ops in a Kubeflow pipeline (or other orchestration tools).
- They are portable across runtime environments.



AI Platform Training with custom containers

```
!gcloud ai-platform jobs submit
training $JOB_ID \
  --region=$REGION \
  --master-image-uri=$IMAGE_URI \
  --training_args
```

Train Models from Images in GCR



AI Platform Training job as an op in a Kubeflow pipeline

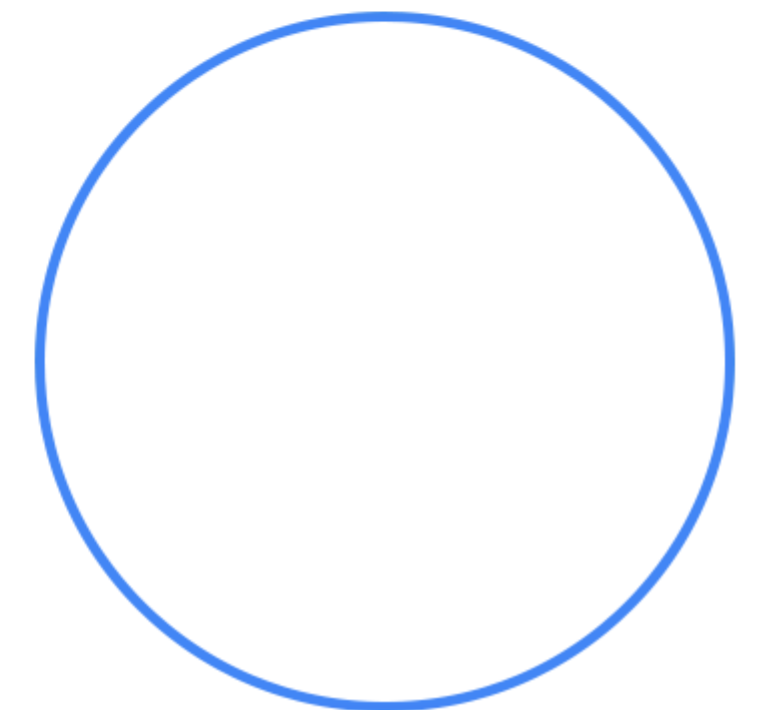
```
import kfp.dsl as dsl

mlengine_train_op =
component_store.load_component(
    'ml_engine/train')

@dsl.pipeline(
    name='My Pipeline'
)
def pipeline(pipeline_args):

    train_model = mlengine_train_op(
        project_id=project_id,
        region=region,
        master_image_uri=TRAINER_IMAGE,
        args=training_args)
```

- Load the pre-built AI Platform Training component.
- Use op with training image and args.



Step 1: Create a model training script

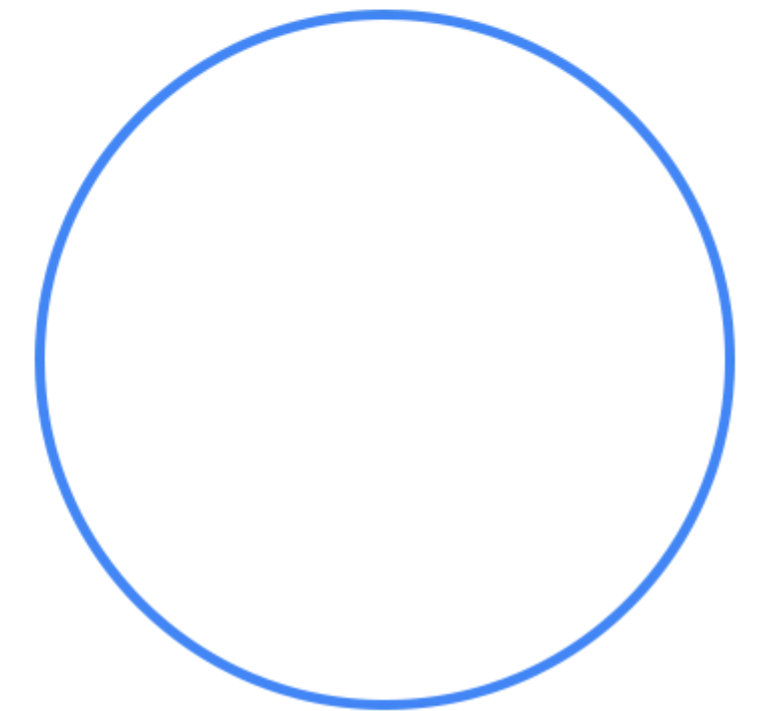
```
%%writefile ./tensorflow_trainer_image/train.py
import tensorflow as tf

def train_evaluate(training_args):
    # Data ingestion and model building code here

    history = model.fit(
        trainds,
        validation_data=evalds,
        epochs=num_evals,
        steps_per_epoch=steps_per_epoch
    )

    tf.saved_model.save(
        obj=model, export_dir=EXPORT_PATH)

if __name__ == '__main__':
    fire.Fire(train_evaluate)
```




Step 2: Create a Dockerfile

```
%%writefile ./trainer_image/Dockerfile

FROM
gcr.io/deeplearning-platform-release/b
ase-cpu
RUN pip install -U fire
tensorflow==2.1.1
WORKDIR /app
COPY train.py .

ENTRYPOINT ["python", "train.py"]
```



📁 / ... / lab / trainer_image /


Name

📄 Dockerfile

🐍 train.py

Step 3: Build the image and push to Container Registry

```
IMAGE_NAME='trainer_image'  
IMAGE_TAG='latest'  
IMAGE_URI=f'gcr.io/{PROJECT_ID}/{IMAGE_NAME}:{IMAGE_TAG}'  
  
!gcloud builds submit --tag $IMAGE_URI $IMAGE_NAME
```



Agenda


Containerized training applications

Containerizing PyTorch, Scikit, and XGBoost applications

Kubeflow and AI Platform Pipelines


Continuous training

What if you want to develop your models with a different framework?



The process is exactly the same!

1. Develop a training script in the framework of your choice.
2. Package the training script into a Docker image.
3. Build and push the image to Container Registry.



PyTorch example

```
%%writefile ./pytorch_trainer_image/train.py
import torch

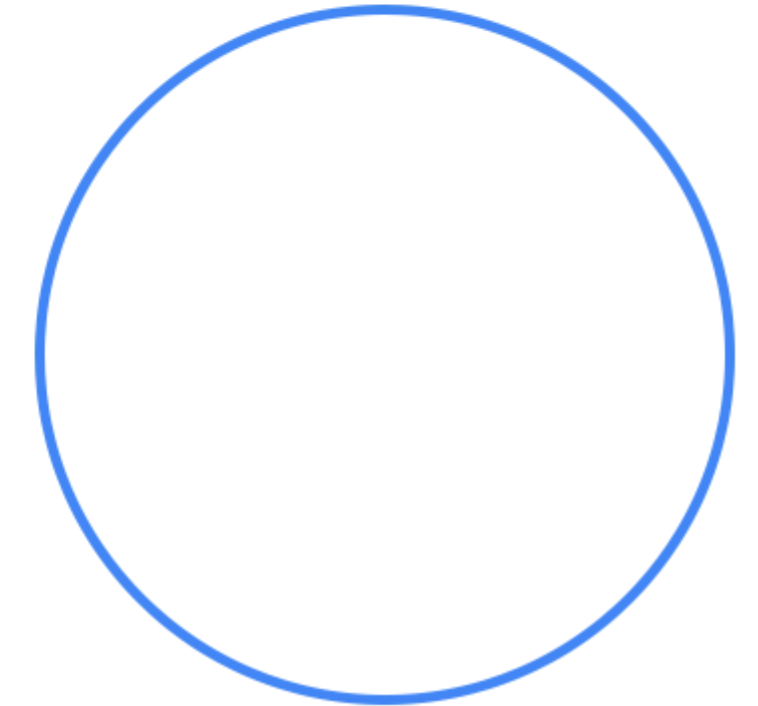
def train_evaluate(training_args):
    # Data ingestion and model building code here
    model.train()
    for e in range(1, num_epochs+1):
        for X_batch, y_batch in train_loader:
            optimizer.zero_grad()
            y_pred = model(X_batch)
            loss = criterion(y_pred, y_batch)
            loss.backward()
            optimizer.step()

    torch.save(model.state_dict(), model_filename)

if __name__ == '__main__':
    fire.Fire(train_evaluate)
```

Step 1:

Develop a training script.




PyTorch example

```
%%writefile ./pytorch_trainer_image/Dockerfile

FROM
gcr.io/deeplearning-platform-release/base-cpu
RUN pip install -U fire torch==1.6.0
scikit-learn==0.23.2 pandas==1.1.1
WORKDIR /app
COPY train.py .

ENTRYPOINT ["python", "train.py"]
```

Step 2:
Create a
Dockerfile.




PyTorch example

```
IMAGE_NAME='pytorch_trainer_image'
IMAGE_TAG='latest'
IMAGE_URI=f'gcr.io/{PROJECT_ID}/{IMAGE_NAME}:{IMAGE_TAG}'

!gcloud builds submit --tag $IMAGE_URI
$IMAGE_NAME
```

Step 3:
Build and push the image.



Agenda


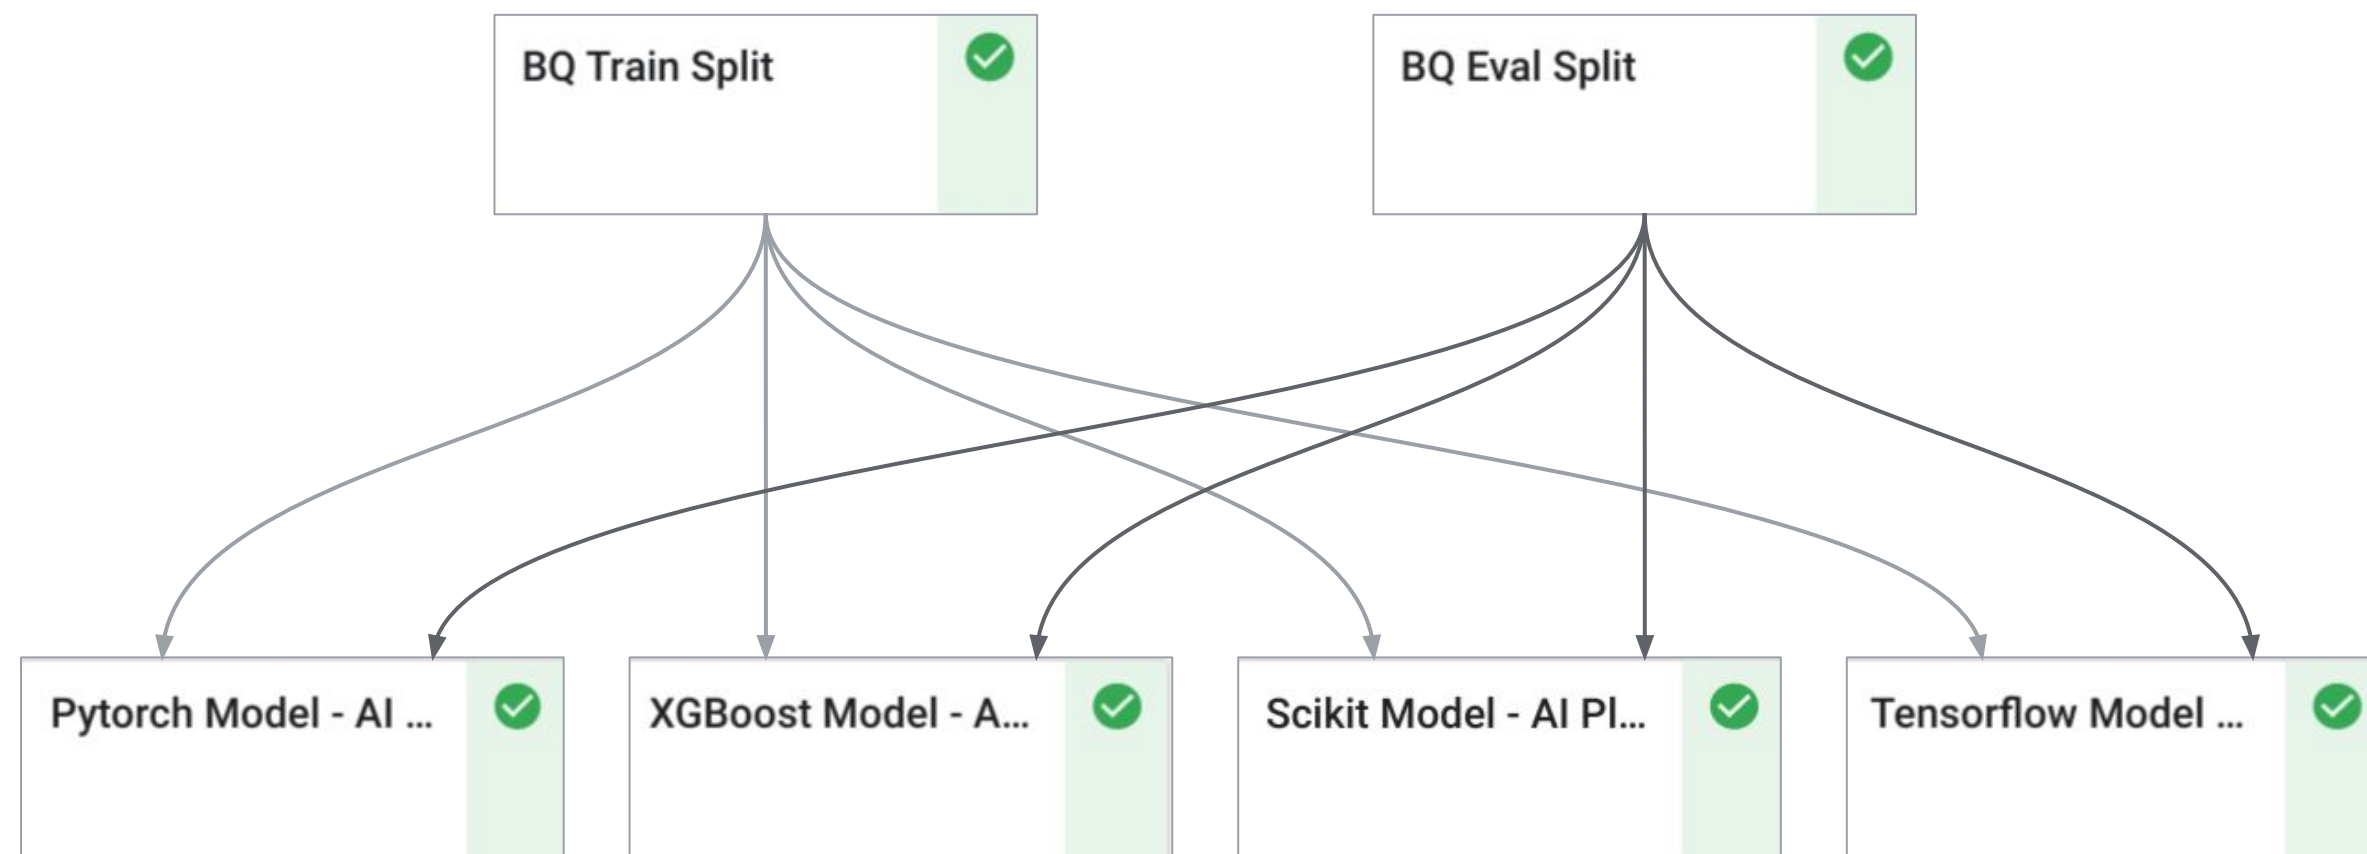
Containerized training applications

Containerizing PyTorch, Scikit, and XGBoost applications

Kubeflow and AI Platform Pipelines


Continuous training

Training multiple models in a Kubeflow pipeline




Create separate arg lists for each container

```
torch_args = [  
    '--training_dataset_path',  
    create_training_split.outputs['output_gcs_path'],  
    '--validation_dataset_path',  
    create_validation_split.outputs['output_gcs_path'],  
    '--output_dir', torch_output_dir,  
    '--batch_size', '32',  
    '--num_epochs', '15',  
]  
  
xgb_args = [  
    '--training_dataset_path',  
    create_training_split.outputs['output_gcs_path'],  
    '--validation_dataset_path',  
    create_validation_split.outputs['output_gcs_path'],  
    '--output_dir', xgb_output_dir,  
    '--max_depth', '10',  
    '--n_estimators', '100',  
]
```



Use multiple mlengine_train ops in pipeline

```
train_torch = mlengine_train_op(  
    project_id=project_id,  
    region=region,  
    master_image_uri=TORCH_TRAINER_IMAGE,  
    args=torch_args).set_display_name('Pytorch  
Model - AI Platform Training')  
  
train_xgb = mlengine_train_op(  
    project_id=project_id,  
    region=region,  
    master_image_uri=XGB_TRAINER_IMAGE,  
    args=xgb_args).set_display_name('XGBoost  
Model - AI Platform Training')
```



Agenda

Containerized training applications



Containerizing PyTorch, Scikit, and XGBoost applications

Kubeflow and AI Platform Pipelines

Continuous training


Considerations for continuous training

Model performance with different retraining intervals




Considerations for continuous training

Model performance with different retraining intervals




- Deterioration of model performance

- Weekly retraining
- Monthly retraining
- Three-month retraining




Considerations for continuous training

Model performance with different retraining intervals




- Deterioration of model performance
- Changes in the data distributions

- Weekly retraining
- Monthly retraining
- Three-month retraining




Considerations for continuous training

Model performance with different retraining intervals



- Deterioration of model performance
- Changes in the data distributions
- Cost and time to retrain model

- Weekly retraining
- Monthly retraining
- Three-month retraining



Scheduled pipeline runs with AI Platform Pipelines

Run Type

☐ One-off

☒ Recurring

Run trigger
Choose a method by which new runs will be triggered

Trigger type*
Periodic

Maximum concurrent runs*
1

☐ Has start date

☐ Has end date

☒ Catchup ?

Run every

7

Days

☐ Has end date

☒ Catchup ?

Run every

7

Minutes

Hours


Days

Weeks

Months

Run parameters
Specify parameters required for the pipeline

train_query



Lab

Continuous Training with TensorFlow, PyTorch, XGBoost, and Scikit Learn Models with Kubeflow and AI Platform Pipelines

In this lab, you create containerized training applications for ML models in multiple frameworks. You will use these images as ops in a Kubeflow pipeline and train them in parallel. You will then set up recurring runs of your Kubeflow pipeline in the UI.

